

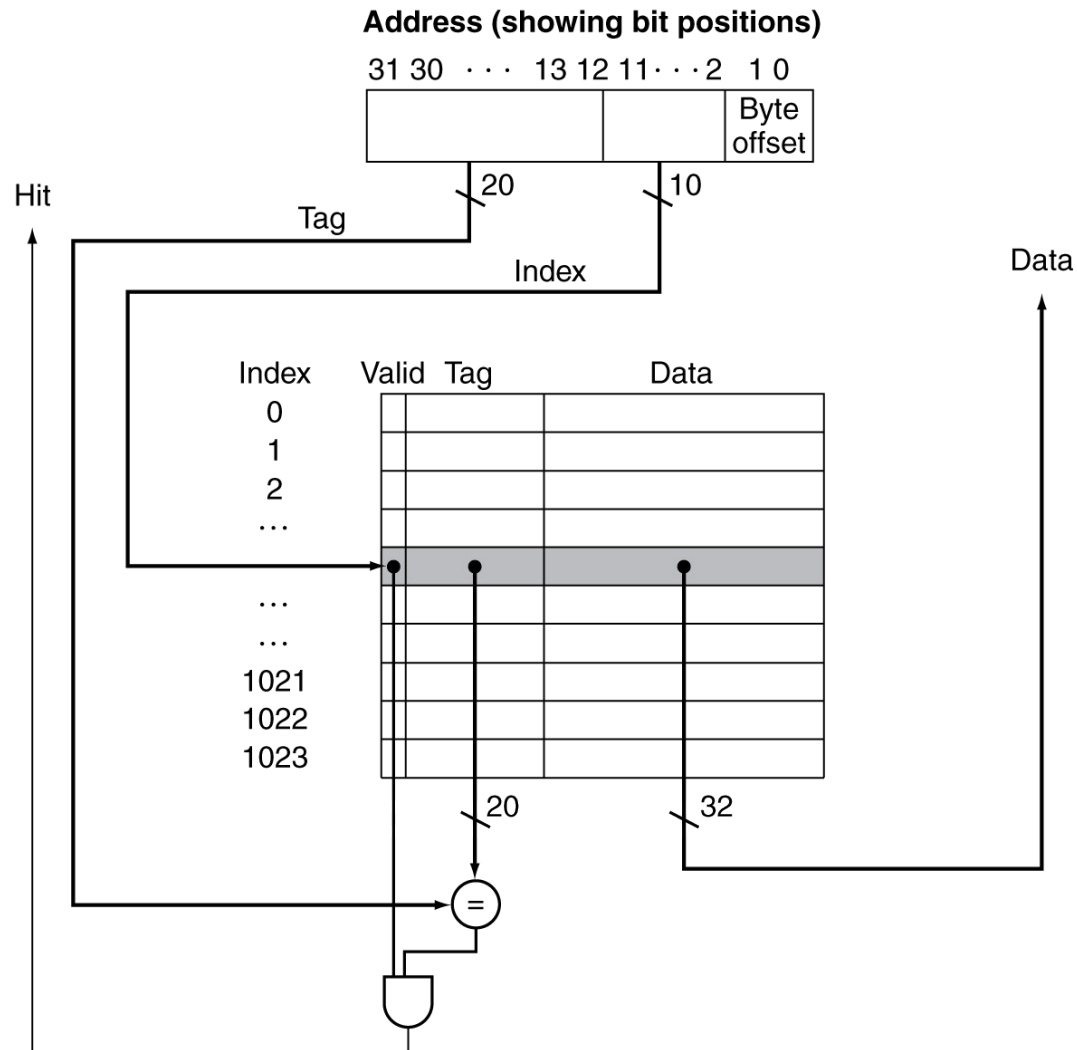
## **Memory System**

**IZTECH, Fall 2023**


**28 December 2023**



# Direct-Mapped Cache - 4KB



# 4-way Set Associative Cache




# Cache Bits in Direct Mapped Cache

## A cache with $2^k$ blocks, each containing $2^n$ bytes

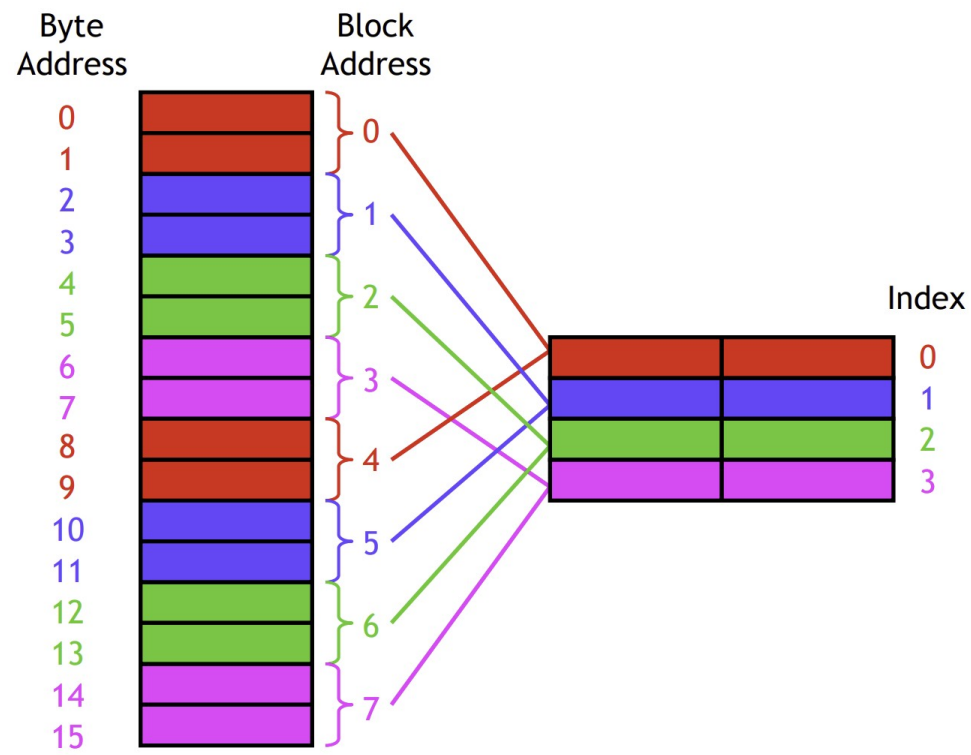
$k$  bits of the address will select one of the  $2^k$  cache blocks

The lowest  $n$  bits are now a block offset that decides which of the  $2^n$  bytes in the cache block will store the data



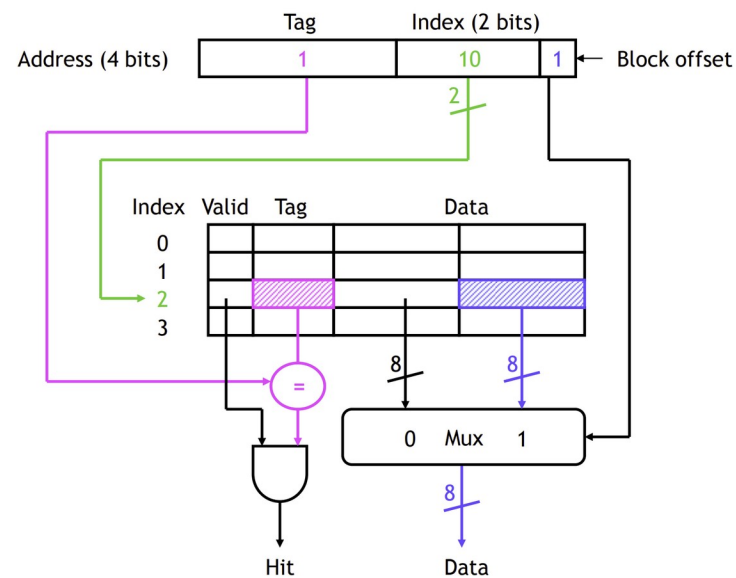

# Example Direct-Mapped Cache

$2^2$  block cache with  $2^1$  bytes per block



# Example Memory Address

## Memory address 13 (1101)



# Multiword Cache Block


**64 blocks, 16 bytes/block**

**To what block number does address 1200 map?**

**Block address =  $1200/16 = 75$**


1 001011 0000

**Block number =  $75 \text{ modulo } 64 = 11$**



# Cache Bits in Associative Caches

**A cache with  $2^s$  sets and each block has  $2^n$  bytes**






# Example Set Associative Cache

**8-block cache with 16 bytes per block**

**Memory address 6195: 00...0110000 011 0011**


1-way associativity  
8 sets, 1 block each



2-way associativity  
4 sets, 2 blocks each



4-way associativity  
2 sets, 4 blocks each



# Cache Operations

## When CPU needs to read data

checks the level 1 cache, then the level 2, so on

if data is available: cache **hit**

otherwise: cache **miss**, accesses the main memory

## When CPU writes data to cache

the line is written to main memory when it is written to the cache: **write-through**

the data isn't written, marked as dirty, when the cache line is replaced by a new cache line from memory, the dirty line is written to memory: **write-back**

# Cache Operations

## When CPU needs to read data

checks the level 1 cache, then the level 2, so on

if data is available: cache **hit**

otherwise: cache **miss**, accesses the main memory

## When CPU writes data to cache

the data isn't written, marked as dirty, when the cache line is replaced by a new cache line from memory, the dirty line is written to memory: **write-back**

the line is written to main memory when it is written to the cache: **write-through**

# When do we write the modified data in a cache to the next level?

**Write-back \*\*: When the block is evicted**

**+ Can combine multiple writes to the same block before eviction**

Potentially saves bandwidth between cache levels + saves energy

**- Need a bit in the tag store indicating the block is “dirty/modified”**

**Write-through: At the time the write happens**

**+ Simpler**

**+ All levels are up to date (Consistency)**

**- More bandwidth intensive; no combining of writes**

# Eviction/Replacement Policy

## Which block in the set to replace on a cache miss?

Any invalid block first

If all are valid: Random, FIFO, Least recently used, Not most recently used, Least frequently used, Least costly to re-fetch, Hybrid replacement policies

Optimal replacement policy?

# LRU

**Evict the least recently accessed block**

**Need to keep track of access ordering of blocks**

**Instead of having true LRU, there are LRU approximations in modern processors**

Not MRU (not most recently used)

## **LRU vs Random**

Average hit rate of LRU and Random is similar

# Manual vs Automatic Management

**Manual: Programmer manages data movement across levels**

- **too painful for programmers on substantial programs**

In some embedded processors (on-chip scratchpad SRAM) and GPUs (called “shared memory”)

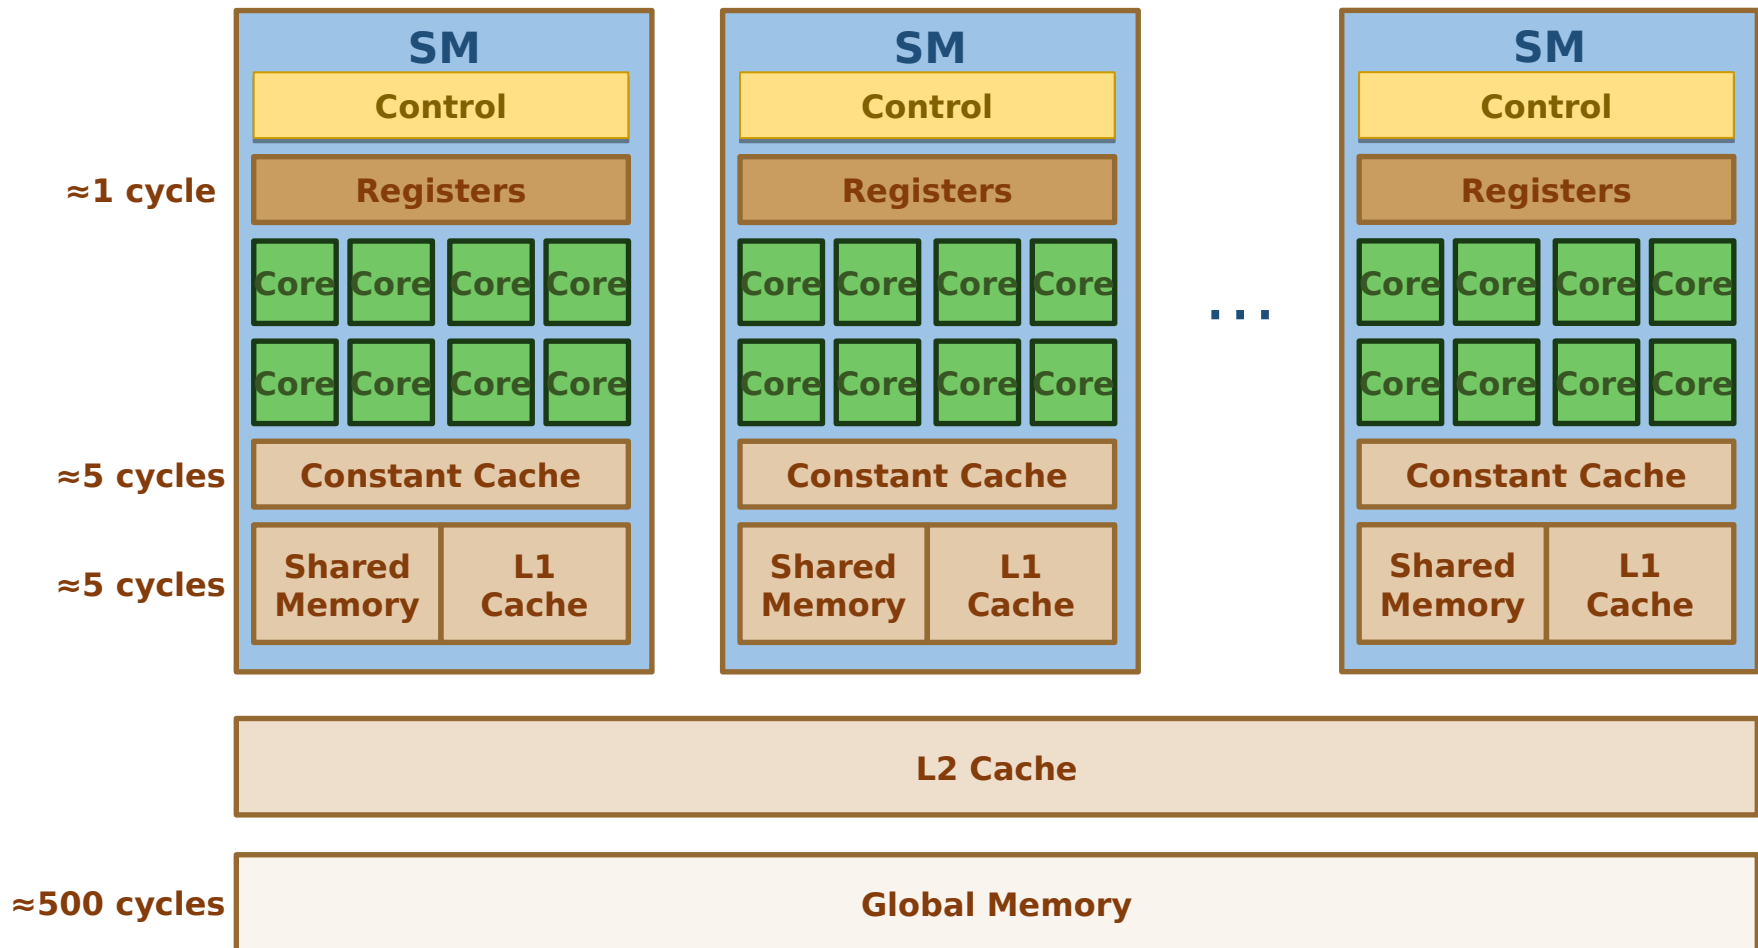
**Automatic: Hardware manages data movement across levels, transparently to the programmer**

- + **programmer’s life is easier**

Average programmer doesn’t need to know about it

You don’t need to know how big the cache is and how it works to write a “correct” program! (What if you want a “fast” program?)

# GPU Memory Structure





# Instruction vs Data Caches

## Unified cache

- + Dynamic sharing of cache space: no overprovisioning that might happen with static partitioning (i.e., separate I and D caches)
- Instructions and data can thrash each other (i.e., no guaranteed space for either)
- I and D are accessed in different places in the pipeline. Where do we place the unified cache for fast access?

**First level caches are almost always split**

**Higher level caches are almost always unified**

# Cache Misses

## Compulsory miss

First reference to an address (block) always results in a miss

Subsequent references should hit unless the cache block is replaced

## Capacity miss

Cache is too small to hold everything needed

Defined as the misses that would occur even in a fully-associative cache (with optimal replacement) of the same capacity

## Conflict miss

Defined as any miss that is neither a compulsory nor a capacity miss

# How to Reduce Misses

## **Compulsory**

Caching cannot help

Prefetching can: Anticipate which blocks will be needed soon

## **Capacity**

Utilize cache space better: keep blocks that will be referenced

Software management: divide working set and computation such that each “computation phase” fits in cache

## **Conflict**

More associativity

# How to Improve Cache Performance

## **Reducing miss rate**

More associativity

Better replacement/insertion policies

## **Reducing miss latency or miss cost**

Multi-level caches

Better replacement/insertion policies

Multiple accesses per cycle

## **Reducing hit latency or hit cost**

# References

## **Chapter 5.3 - 5.4**

**(Computer Organization and Design: The Hardware/Software Interface by Hennessy/Patterson, 5th edition)**